

```
1 !apt-get --purge remove cuda nvidia* libnvidia-*
2 !dpkg -l | grep cuda- | awk '{print $2}' | xargs -n1 dpkg --purge
3 !apt-get remove cuda-*
4 !apt autoremove
5 !apt-get update
```

```
1 !wget https://developer.nvidia.com/compute/cuda/9.2/Prod/local_installers/cuda-repo-ubu
2 !dpkg -i cuda-repo-ubuntu1604-9-2-local_9.2.88-1_amd64.deb
3 !apt-key add /var/cuda-repo-9-2-local/7fa2af80.pub
4 !apt-get update
5 !apt-get install cuda-9.2
```

```
1 !nvcc --version
```

```
1 !pip install git+git://github.com/andreinechaev/nvcc4jupyter.git
```

```
1 %load_ext nvcc_plugin
```

```
1 %%cu
2
3 #include<iostream>
4 #include<math.h>
5 using namespace std;
6
7 __global__
8 void matrixMultiplication(int *a, int *b, int *c, int m, int n, int k)
9 {
10     int row = blockIdx.y*blockDim.y + threadIdx.y;
11     int col = blockIdx.x*blockDim.x + threadIdx.x;
12     int sum=0;
13
14     if(col<k && row<m) {
15         for(int j=0;j<n;j++)
16         {
17             //Converting matrices to 1D array
18             sum += a[row*n+j] * b[j*k+col];
19         }
20         c[k*row+col]=sum;
21     }
22
23 }
24
25 void init_result(int *a, int m, int k) {
26     for(int i=0; i<m; i++) {
27         for(int j=0; j<k; j++) {
28             a[i*k + j] = 0;
29         }
30     }
31 }
32
```

```

33 void init_matrix(int *a, int n, int m) {
34     for(int i=0; i<n; i++) {
35         for(int j=0; j<m; j++) {
36             a[i*m + j] = rand()%10 + 1;
37         }
38     }
39 }
40
41 void print_matrix(int *a, int n, int m) {
42     for(int i=0; i<n; i++) {
43         for(int j=0; j<m; j++) {
44             cout<<" "<<a[i*m + j];
45         }
46         cout<<endl;
47     }
48     cout<<endl;
49 }
50
51 int main()
52 {
53     cout<<"Matrix Multiplication: "<<endl;
54     int *a,*b,*c;
55     int *a_dev,*b_dev,*c_dev;
56     int m=5, n=4, k=3;
57
58     a = new int[m*n];
59     b = new int[n*k];
60     c = new int[m*k];
61
62     init_matrix(a, m, n);
63     init_matrix(b, n ,k);
64     init_result(c, m, k);
65
66     cout<<"Initial matrix : "<<endl;
67
68     print_matrix(a, m, n);
69     print_matrix(b, n, k);
70
71     cudaMalloc(&a_dev, sizeof(int)*m*n);
72     cudaMalloc(&b_dev, sizeof(int)*n*k);
73     cudaMalloc(&c_dev, sizeof(int)*m*k);
74
75     cudaMemcpy(a_dev, a, sizeof(int)*m*n, cudaMemcpyHostToDevice);
76     cudaMemcpy(b_dev, b, sizeof(int)*n*k, cudaMemcpyHostToDevice);
77
78     //Defining dimensions
79     dim3 dimGrid(1,1);
80     dim3 dimBlock(16,16);
81     matrixMultiplication<<<dimGrid, dimBlock>>>(a_dev,b_dev,c_dev, m, n, k);
82
83     cudaMemcpy(c, c_dev, sizeof(int)*m*k, cudaMemcpyDeviceToHost);
84
85     cout<<"Result : "<<endl;
86     print_matrix(c, m, k);
87
88     . . .

```

```

88     cudaFree(a_dev);
89     cudaFree(b_dev);
90     cudaFree(c_dev);
91
92     delete[] a;
93     delete[] b;
94     delete[] c;
95     return 0;
96 }

```

☞ Matrix Multiplication:

Initial matrix :

```

4  7  8  6
4  6  7  3
10 2  3  8
1  10 4  7
1  7  3  7

```

```

2  9  8
10 3  1
3  4  8
6  10 3

```

Result :

```

138 149 121
107 112 103
97  188 130
156 125 71
123 112 60

```

```

1 %%cu
2
3 #include<iostream>
4 #include<math.h>
5 using namespace std;
6
7 __global__
8 void matrixVector(int *vec, int *mat, int *result, int n, int m)
9 {
10     int tid = blockIdx.x*blockDim.x + threadIdx.x;
11     int sum=0;
12
13     if(tid <= n) {
14         for(int i=0; i<n; i++) {
15             sum += vec[i]*mat[(i*m) + tid];
16         }
17         result[tid] = sum;
18     }
19 }
20
21 void init_array(int *a, int n) {
22     for(int i=0; i<n; i++)
23         a[i] = rand()%n + 1;
24 }
25

```

```

26 void init_matrix(int *a, int n, int m) {
27     for(int i=0; i<n; i++) {
28         for(int j=0; j<m; j++) {
29             a[i*m + j] = rand()%n + 1;
30         }
31     }
32 }
33
34 void print_array(int *a, int n) {
35     for(int i=0; i<n; i++) {
36         cout<<" "<<a[i];
37     }
38     cout<<endl;
39 }
40
41 void print_matrix(int *a, int n, int m) {
42     for(int i=0; i<n; i++) {
43         for(int j=0; j<m; j++)
44             cout<<" "<<a[i*m + j];
45         cout<<endl;
46     }
47 }
48
49 int main() {
50     cout<<"Vector and Matrix Multiplication: "<<endl;
51     int *a, *b, *c;
52     int *a_dev, *b_dev, *c_dev;
53
54     int n = 3;
55     int m = 4;
56
57     a = new int[n];
58     b = new int[n*m];
59     c = new int[m];
60
61     init_array(a, n);
62     init_matrix(b, n, m);
63
64     cout<<"Initial array : "<<endl;
65     print_array(a, n);
66     cout<<"Initial matrix : "<<endl;
67     print_matrix(b, n, m);
68     cout<<"Initial resultant array : "<<endl;
69     print_array(c, m);
70     cout<<endl;
71
72     cudaMalloc(&a_dev, sizeof(int)*n);
73     cudaMalloc(&b_dev, sizeof(int)*n*m);
74     cudaMalloc(&c_dev, sizeof(int)*m);
75
76     cudaMemcpy(a_dev, a, sizeof(int)*n, cudaMemcpyHostToDevice);
77     cudaMemcpy(b_dev, b, sizeof(int)*n*m, cudaMemcpyHostToDevice);
78
79     matrixVector<<<m/256+1, 256>>>(a_dev, b_dev, c_dev, n, m);
80

```

```

81     cudaMemcpy(c, c_dev, sizeof(int)*m, cudaMemcpyDeviceToHost);
82
83     cout<<"Result : "<<endl;
84     print_array(c, m);
85
86     cudaFree(a_dev);
87     cudaFree(b_dev);
88     cudaFree(c_dev);
89
90     delete[] a;
91     delete[] b;
92     delete[] c;
93
94     return 0;
95 }

```

☞ Vector and Matrix Multiplication:

Initial array :

2 2 1

Initial matrix :

2 3 2 2

1 1 2 3

2 3 2 3

Initial resultant array :

0 0 0 0

Result :

8 11 10 13

```

1 %%cu
2
3 #include<iostream>
4 #include<math.h>
5 using namespace std;
6
7 __global__
8 void add(int *a, int *b, int *result, int n) {
9     int index = blockIdx.x*blockDim.x + threadIdx.x;
10    if(index <= n) {
11        result[index] = a[index] + b[index];
12    }
13 }
14
15 void print_array(int *a, int N) {
16     for(int i=0; i<N; i++) {
17         cout<<" "<<a[i];
18     }
19     cout<<endl;
20 }
21
22 void init_array(int *a, int N) {
23     for(int i=0; i<N; i++) {
24         a[i] = rand()%100 + 1;
25     }
26 }

```

```

27
28 int main() {
29     cout<<"Addition of two large vectors : "<<endl;
30     int *a, *b, *c;
31     int *a_dev, *b_dev, *c_dev;
32     int n = 100;
33     int threads_per_block = 25;
34     int size = n * sizeof(int);
35
36     a = new int[n];
37     b = new int[n];
38     c = new int[n];
39
40     init_array(a, n);
41     init_array(b, n);
42
43     cudaMalloc(&a_dev, size);
44     cudaMalloc(&b_dev, size);
45     cudaMalloc(&c_dev, size);
46
47     cout<<"Array 1 : "<<endl;
48     print_array(a, n);
49     cout<<"Array 2 : "<<endl;
50     print_array(b, n);
51
52     cudaMemcpy(a_dev, a, size, cudaMemcpyHostToDevice);
53     cudaMemcpy(b_dev, b, size, cudaMemcpyHostToDevice);
54     add<<<n/threads_per_block, threads_per_block>>>(a_dev, b_dev, c_dev, n);
55     cudaMemcpy(c, c_dev, size, cudaMemcpyDeviceToHost);
56
57     cout<<"Result : "<<endl;
58     print_array(c, n);
59
60     cudaFree(a_dev);
61     cudaFree(b_dev);
62     cudaFree(c_dev);
63
64     return 0;
65 }

```

➡ Addition of two large vectors :

Array 1 :

84 87 78 16 94 36 87 93 50 22 63 28 91 60 64 27 41 27 73 37 12

Array 2 :

96 71 35 79 68 2 98 3 18 93 53 57 2 81 87 42 66 90 45 20 41 36

Result :

180 158 113 95 162 38 185 96 68 115 116 85 93 141 151 69 107 117